

# Building Machine Learning Models in Spark 2

---

MACHINE LEARNING PACKAGES: SPARK.MLLIB VS. SPARK.ML



**Janani Ravi**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**Spark 1.x provided powerful support for ML in spark.mllib**

**Spark 2.x goes further with spark.ml**


**Faster execution**

**Ease of hyperparameter tuning**

**Both libraries offer powerful support**

# Prerequisites and Course Outline

---



# Prerequisite Courses - Spark


## **Beginning Data Exploration and Analysis with Apache Spark**

- Programming in Spark 1.x using Python

## **Getting Started with Spark 2**

- Programming in Spark 2.x using Python

# Prerequisite Courses - ML



## **How to Think About Machine Learning Algorithms**

- Basic ML understanding and algorithms

## **Understanding Machine Learning with Python**

- Basic ML in the Python language



# Software and Skills

**Be very comfortable programming in Python (Python 3)**

**Be comfortable working with Jupyter notebooks**

**Understand the basics of Spark and ML**



# Course Outline

## **ML libraries in Spark 1 vs Spark 2**

- Basic Spark architecture
- Why 2 libraries, when to use one over the other?

## **Classification and regression models**

- Decisions trees, random forest
- Linear regression, Lasso and Ridge regression

## **Clustering and dimensionality reduction**

- k-means clustering, PCA

## **Recommendation systems**

- Collaborative filtering using explicit and implicit ratings

# RDDs and Spark 1.x


---



Why is this relevant in Spark 2?


RDDs are still the **fundamental building blocks** of Spark

# Resilient Distributed Datasets



**All operations in Spark are performed on *in-memory* objects**

# Resilient Distributed Datasets



**An RDD is a **collection** of entities**  
**- rows, records**

# Characteristics of RDDs

**Partitioned**

**Split across data nodes in a cluster**


**Immutable**

**RDDs, once created, cannot be changed**

**Resilient**

**Can be reconstructed even if a node crashes**

# Partitioned



# Partitions

**RDDs  
represent data  
in-memory**

1	Indigo	06:45	Bangalore
2	Jet Air	08:45	New Delhi
3	SpiceJet	09:15	Mumbai
4	Indigo	10:45	New Delhi
5	Air India	11:15	Mumbai
6	Vistara	12:00	New Delhi

**Data is divided  
into partitions**

## Partitions

1	Indigo	06:45	Bangalore
2	Jet Air	08:45	New Delhi
3	SpiceJet	09:15	Mumbai
4	Indigo	10:45	New Delhi
5	Air India	11:15	Mumbai
6	Vistara	12:00	New Delhi

**Data is divided  
into partitions**

## Partitions

1	Indigo	06:45	Bangalore
2	Jet Air	08:45	New Delhi

3	SpiceJet	09:15	Mumbai
4	Indigo	10:45	New Delhi

5	Air India	11:15	Mumbai
6	Vistara	12:00	New Delhi



**Distributed to  
multiple  
machines,  
called nodes**

## Partitions

1	Indigo	06:45	Bangalore
2	Jet Air	08:45	New Delhi

3	SpiceJet	09:15	Mumbai
4	Indigo	10:45	New Delhi

5	Air India	11:15	Mumbai
6	Vistara	12:00	New Delhi




**Distributed to  
multiple  
machines,  
called nodes**

Partitions




**Nodes  
process data  
in parallel**

Partitions



Immutable



**An RDD cannot be  
mutated**

**Only *two* operations are  
permitted on an RDD**

# Only Two Types of Operations



**Transformation**


**Transform into  
another RDD**

**Action**

**Request a result**


Transformations are **executed**  
only when a result is requested

# Lazy Evaluation



**Spark keeps a record  
of the series of  
transformations  
requested by the user**


# Lazy Evaluation



**It groups the transformations in an efficient way when an Action is requested**



**Resilient**



**RDDs can be  
reconstructed even if  
the node it lives on  
crashes**

RDDs Are Resilient

**RDDs can be created  
in 2 ways**




**Reading a file**



**Transforming  
another RDD**

# RDDs Are Resilient



Reading a file



Transforming  
another RDD



Every RDD keeps track  
of **where** it came from

# RDDs Are Resilient




It tracks **every** transformation  
which led to the current RDD



# RDDs Are Resilient



**However many  
transformations it  
takes**



# RDDs Are Resilient




This is the  
RDD's **lineage**

# RDDs Are Resilient



None of the  
transformations  
are **applied** till we  
**access** the results



# Characteristics of RDDs

**Partitioned**

**Split across data  
nodes in a cluster**

**Immutable**

**RDD once created  
cannot be changed**

**Resilient**

**Can be  
reconstructed even  
if a node crashes**





# RDDs, DataFrames, Datasets

---

# DataFrame: Data in Rows and Columns

DATE	OPEN	...	PRICE
2016-12-01	772	...	779
2016-11-01	758	...	747
2006-01-01	302	...	309

Each row represents  
1 observation

# DataFrame: Data in Rows and Columns

Each column  
represents 1 variable  
(a list or vector)

DATE	OPEN	...	PRICE
2016-12-01	772	...	779
2016-11-01	758	...	747
2006-01-01	302	...	309

# From File to DataFrame

DATE	OPEN	...	PRICE
2016-12-01	772	...	779
2016-11-01	758	...	747
2006-01-01	302	...	309

File



DATE	OPEN	...	PRICE
2016-12-01	772	...	779
2016-11-01	758	...	747
2006-01-01	302	...	309

DataFrame

# RDDs to DataFrames

## **RDDs**

**Primary abstraction since initial versions**

**Immutable and distributed**

**Conceptually similar to a collection of records**

**No concept of columns**

**No optimized execution**

**Available in all languages**

## **DataFrames**

**Added to Spark in 1.3**

**Also immutable and distributed**

**Conceptually equal to a table in an RDBMS**

**Named columns like Pandas or R**

**Leverage optimizers in recent versions**

**Available in all languages**

# Datasets to DataFrames

## Datasets

**Scala and Java\***

**Type safe OOP interface**

\*Datasets of the Row() object in Scala/  
Java often called DataFrames

## DataFrames

**Python, R, Scala, Java**

**No type safety at compile time**

Equivalent to Dataset<Row> in Java or  
Dataset[Row] in Scala

Starting Spark 2.0, APIs for  
Datasets and DataFrames  
have merged



# DataFrames Built on Top of RDDs

**Partitioned**

**Split across data  
nodes in a cluster**

**Immutable**

**Once created,  
cannot be changed**

**Resilient**

**Can be  
reconstructed even  
if a node crashes**

# Demo

**Install standalone Spark on your local machine**

**Set up the PySpark REPL interface**

# Making the Choice Between spark.ml vs. spark.mllib

---

# Changes Starting Spark 2.0



**Easier**

Unifying Datasets and  
DataFrames, SQL support...



**Faster**

Optimize like a compiler, not a  
DBMS

# Performance Improvements

Comparison of time per row, on 1 billion records on single thread

Primitive	Spark 1.6	Spark 2.0	Speedup Factor
filter	15ns	1.1ns	13.6
sum w/o group	14ns	0.9ns	15.6
sum w/ group	79ns	10.7ns	7.4
hash join	115ns	4.0ns	28.8
sort (8-bit)	620ns	5.3ns	117.0
sort (64-bit)	620ns	40ns	15.5
sort-merge-join	750ns	700ns	1.1

Source: <https://databricks.com/blog/2016/07/26/introducing-apache-spark-2-0.html>



## Ease of Use

**Unified API for DataFrames**

**spark.ml and ML pipelines**

**Advanced streaming**

# spark.mllib and spark.ml

## **spark.mllib**

**Older**

**RDDs**

**For now, more functionality**

**ETL hard - no pipeline support**

**Hyperparameter tuning hard**

## **spark.ml**

**Newer**

**DataFrames (faster!)**

**Functionality catching up**

**Support for ML pipelines**

**Tools for hyperparameter tuning**

# spark.mllib and spark.ml

## spark.mllib

To maintain backward compatibility  
with 1.x applications

To use features which are not available  
in the newer version

ETL is not important, do not need  
pipelines

## spark.ml

Spark 2 is available and you want to  
take advantage of better performance

To use higher levels APIs and  
abstractions for faster development

ETL, chaining transformations  
significant



Both packages are currently useful

- **spark.mllib** has more features

**spark.ml** - feature compatibility  
around the corner

**spark.mllib** will be deprecated in  
the future

# Decision Trees

---

# Jockey or Basketball Player?



## **Jockeys**

**Tend to be light to meet horse carrying limits**



## **Basketball Players**

**Tend to be tall, strong and heavy**



# Jockey or Basketball Player?

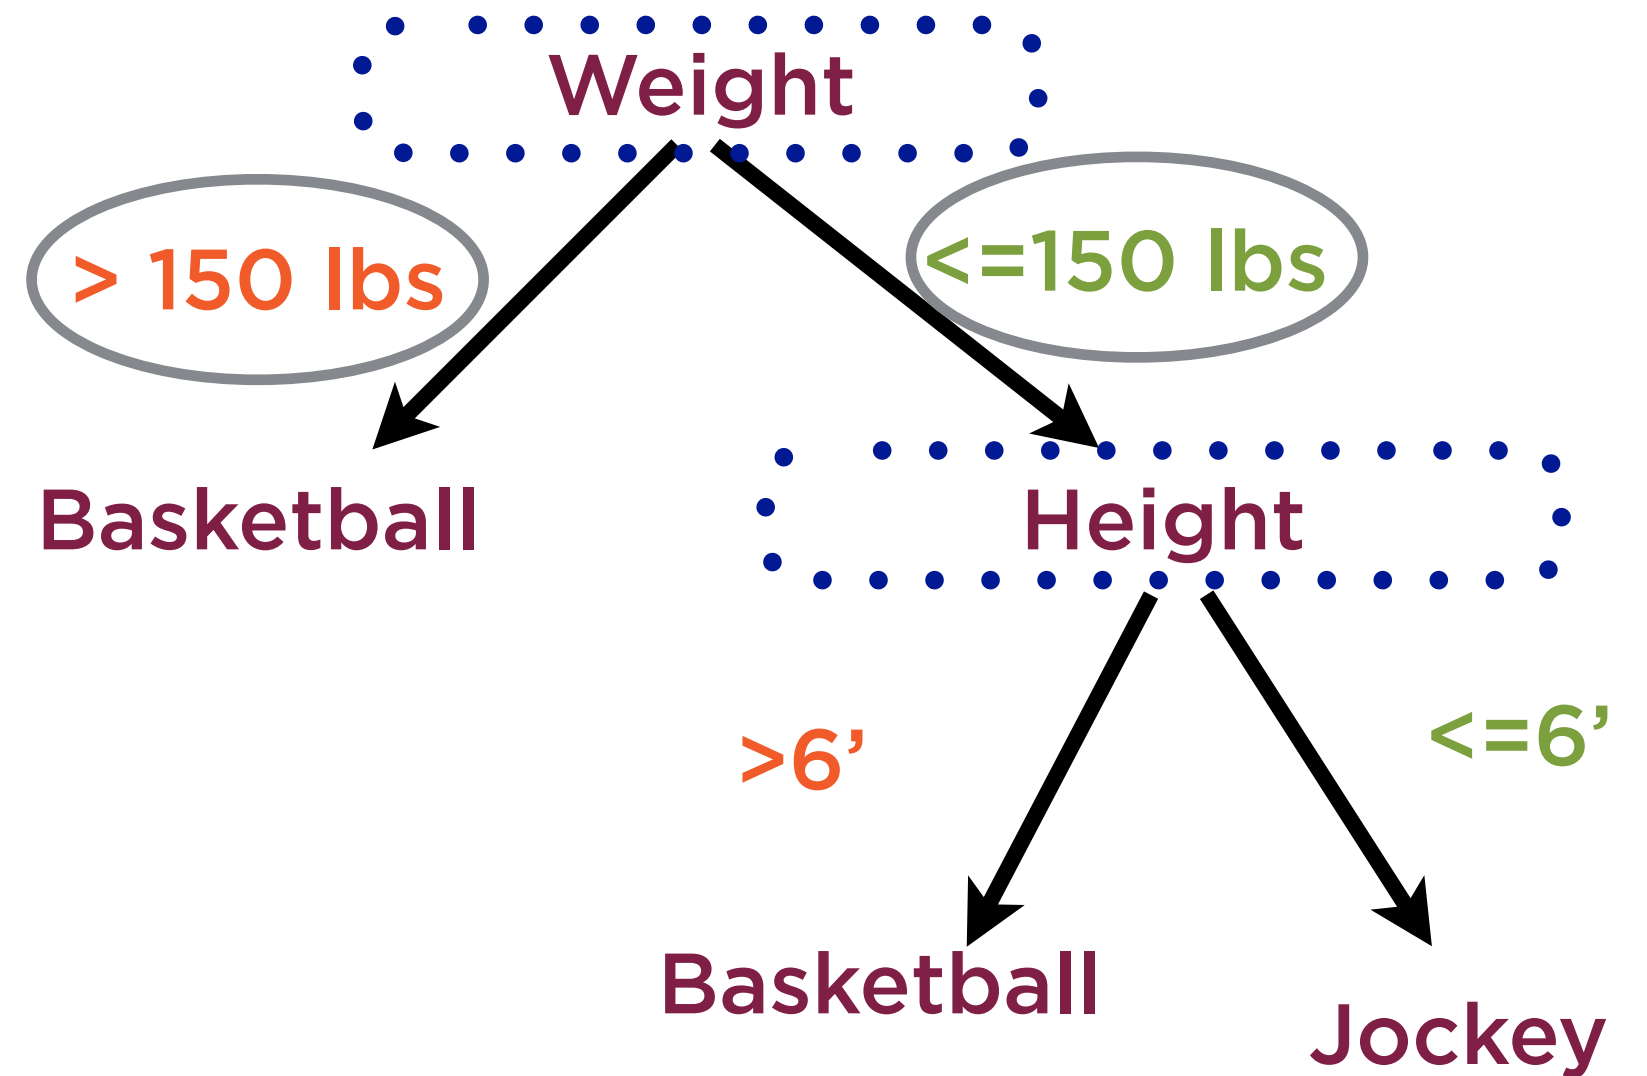
**Intuitively know**

- **jockeys tend to be light...**
- **...and not very tall**
- **basketball players tend to be tall**
- **...and also quite heavy**

Fit knowledge  
into rules

Each rule involves  
a threshold

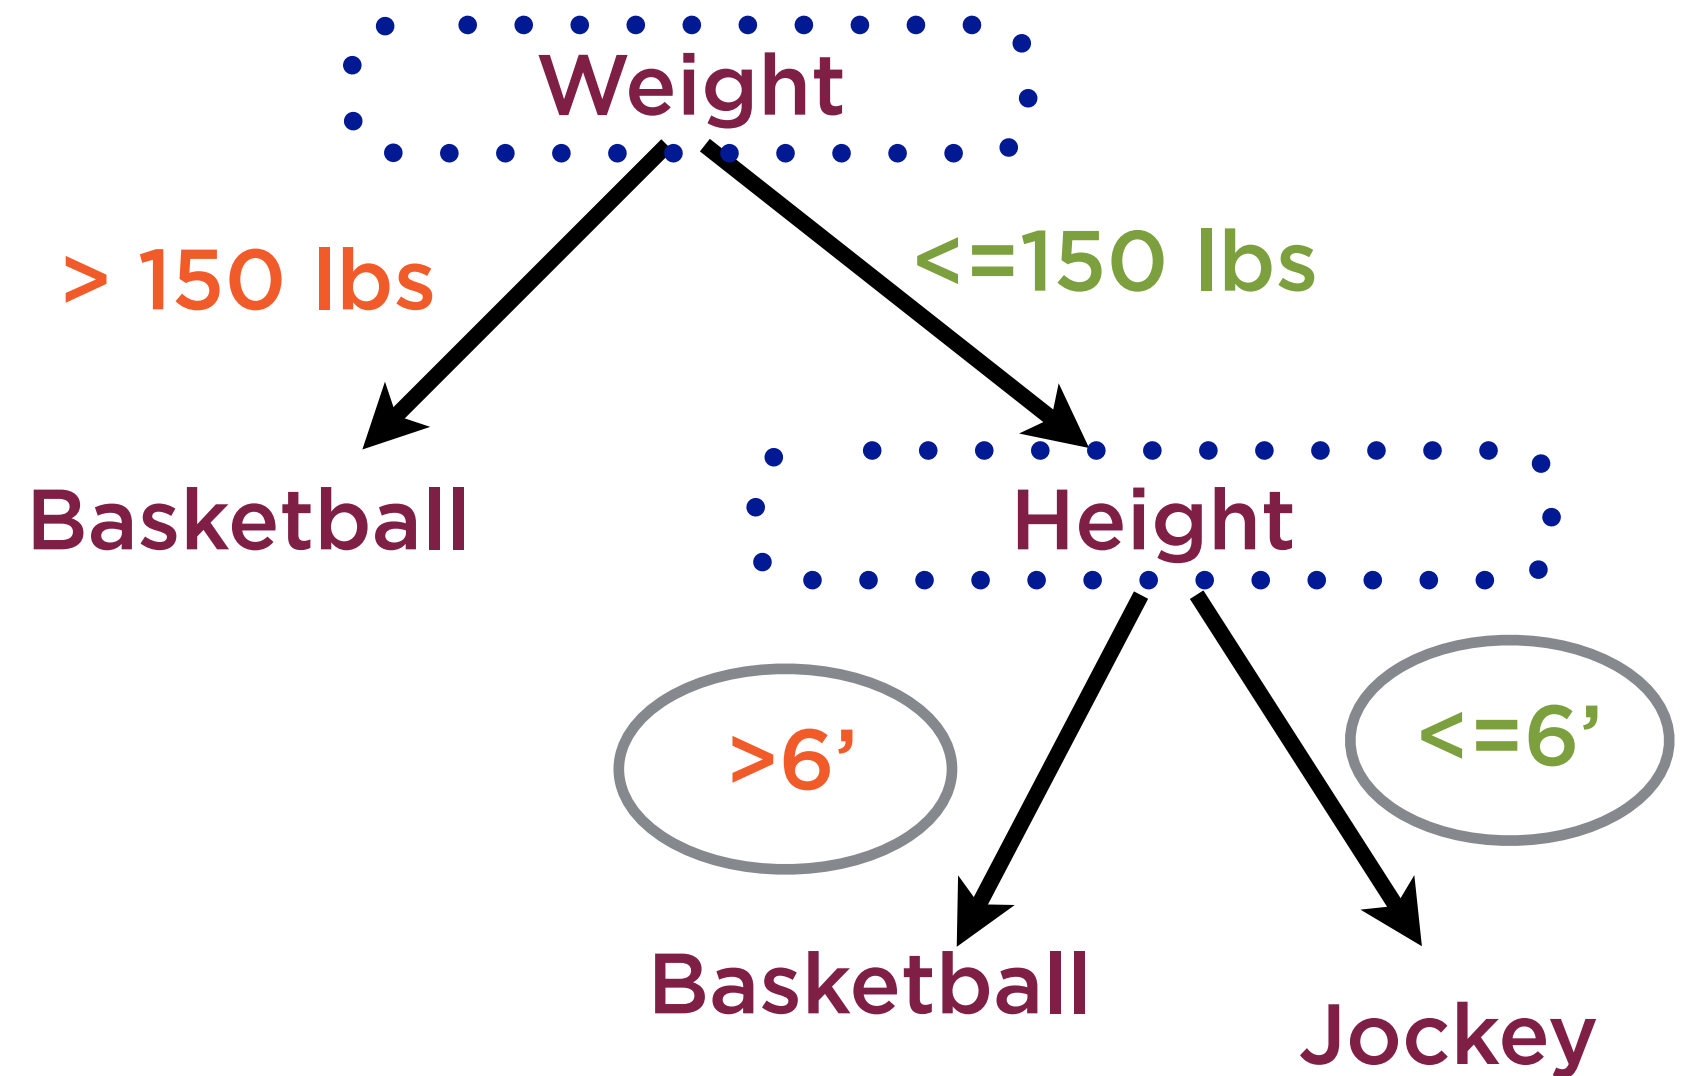
# Decision Tree



Fit knowledge  
into rules

Each rule involves  
a threshold


# Decision Tree



Order of decision  
variables matters

Rules and order  
found using ML

# Decision Tree






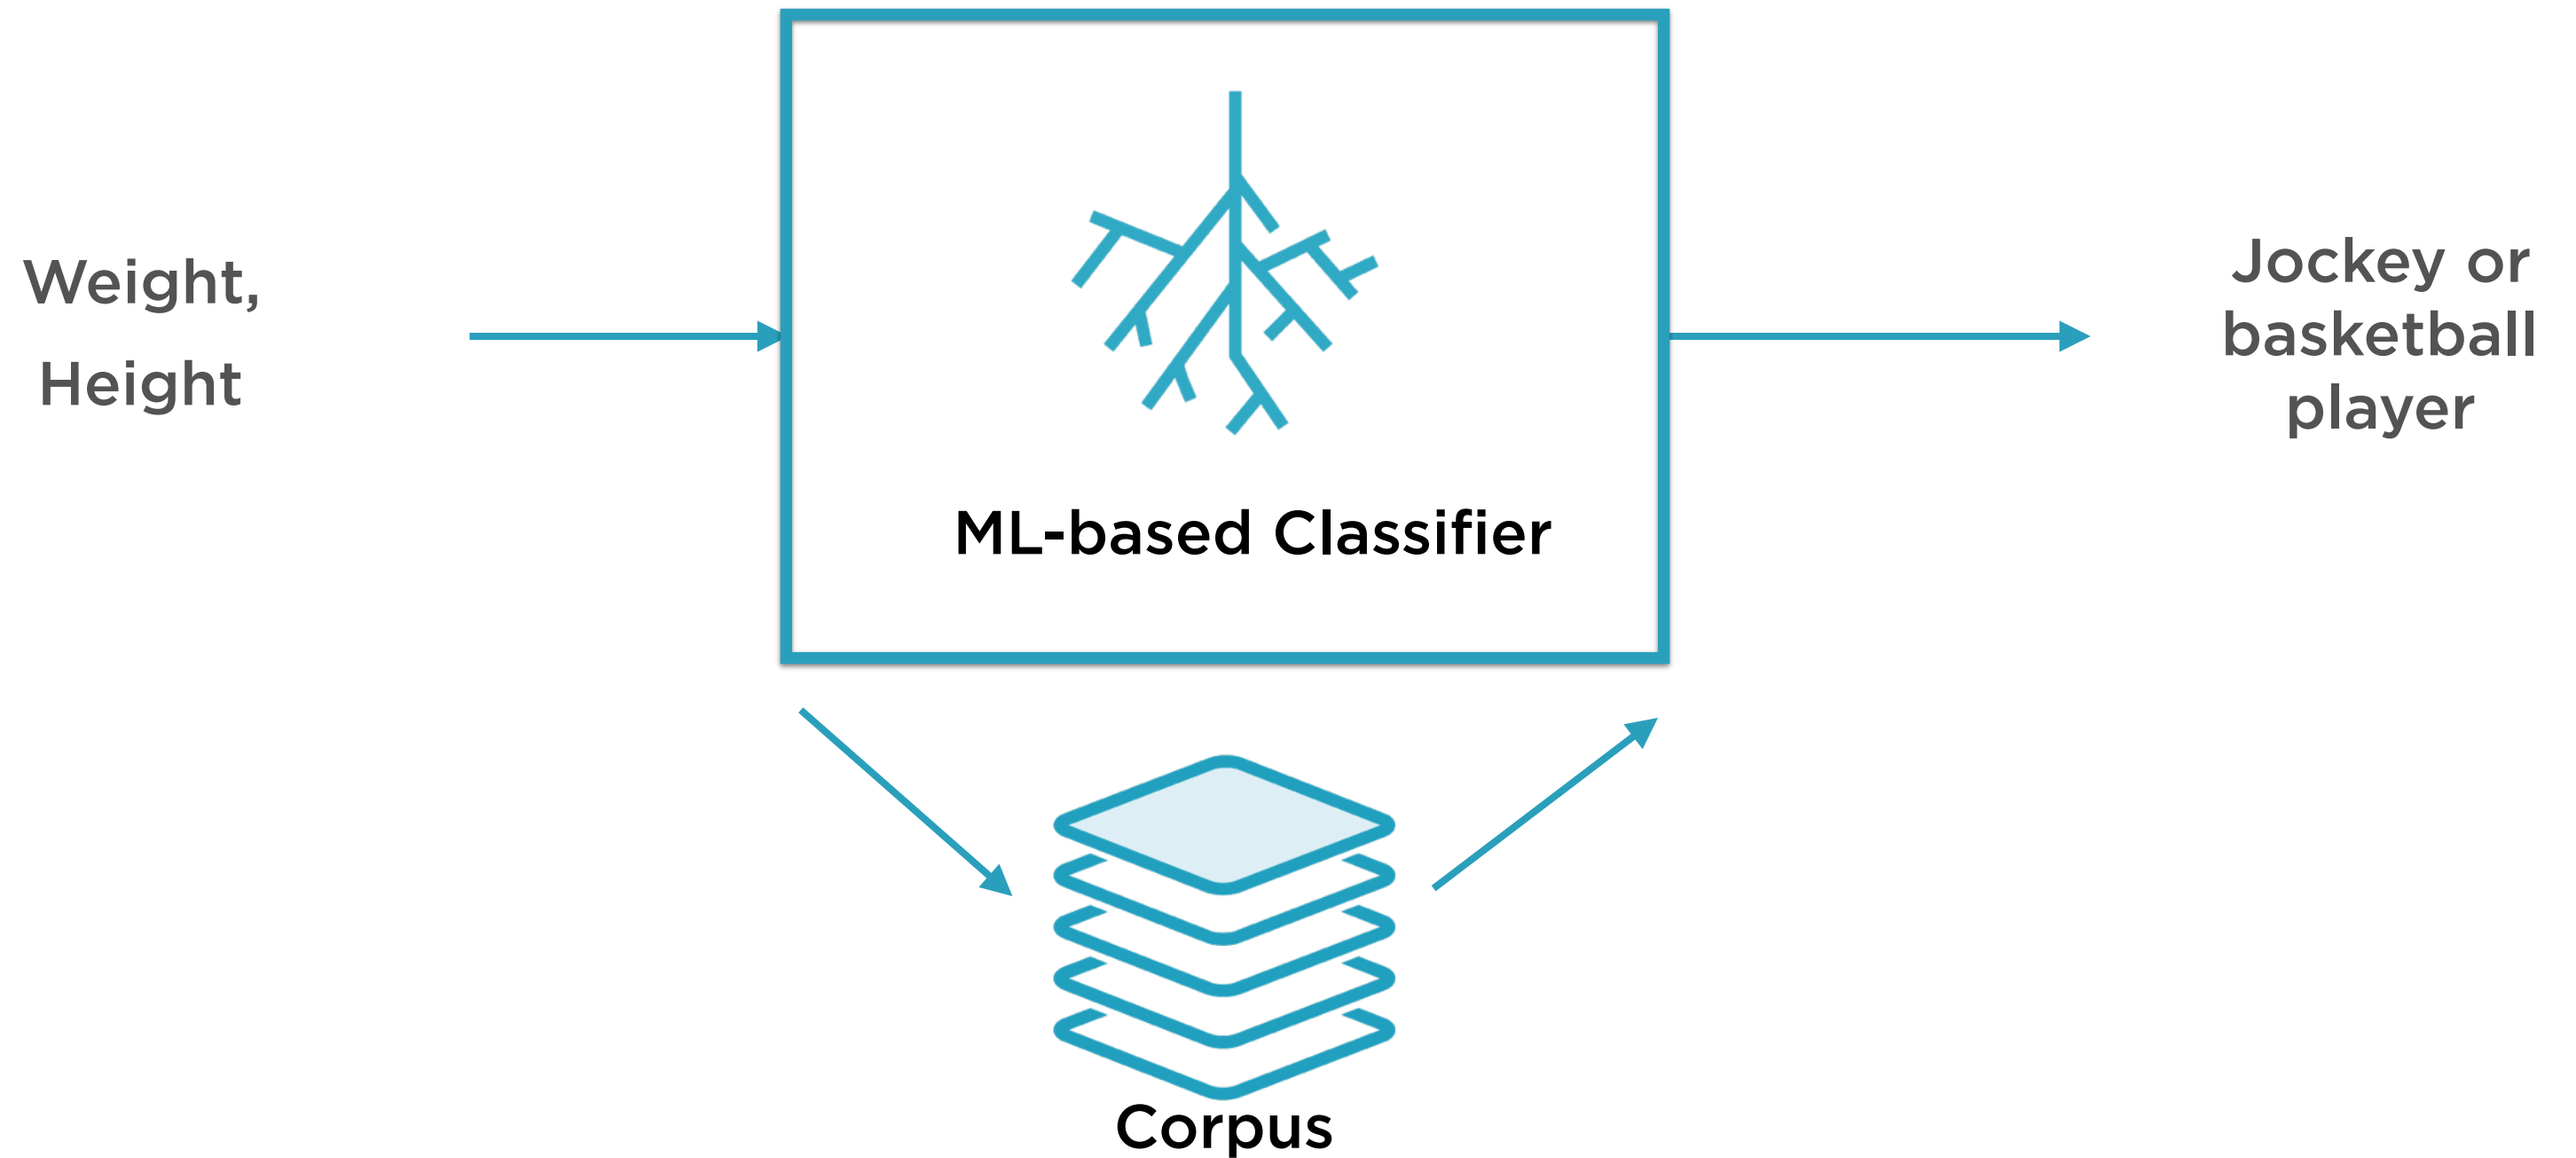
Classification And  
Regression Tree

“CART”


# Decision Tree



# Decision Trees for Classification



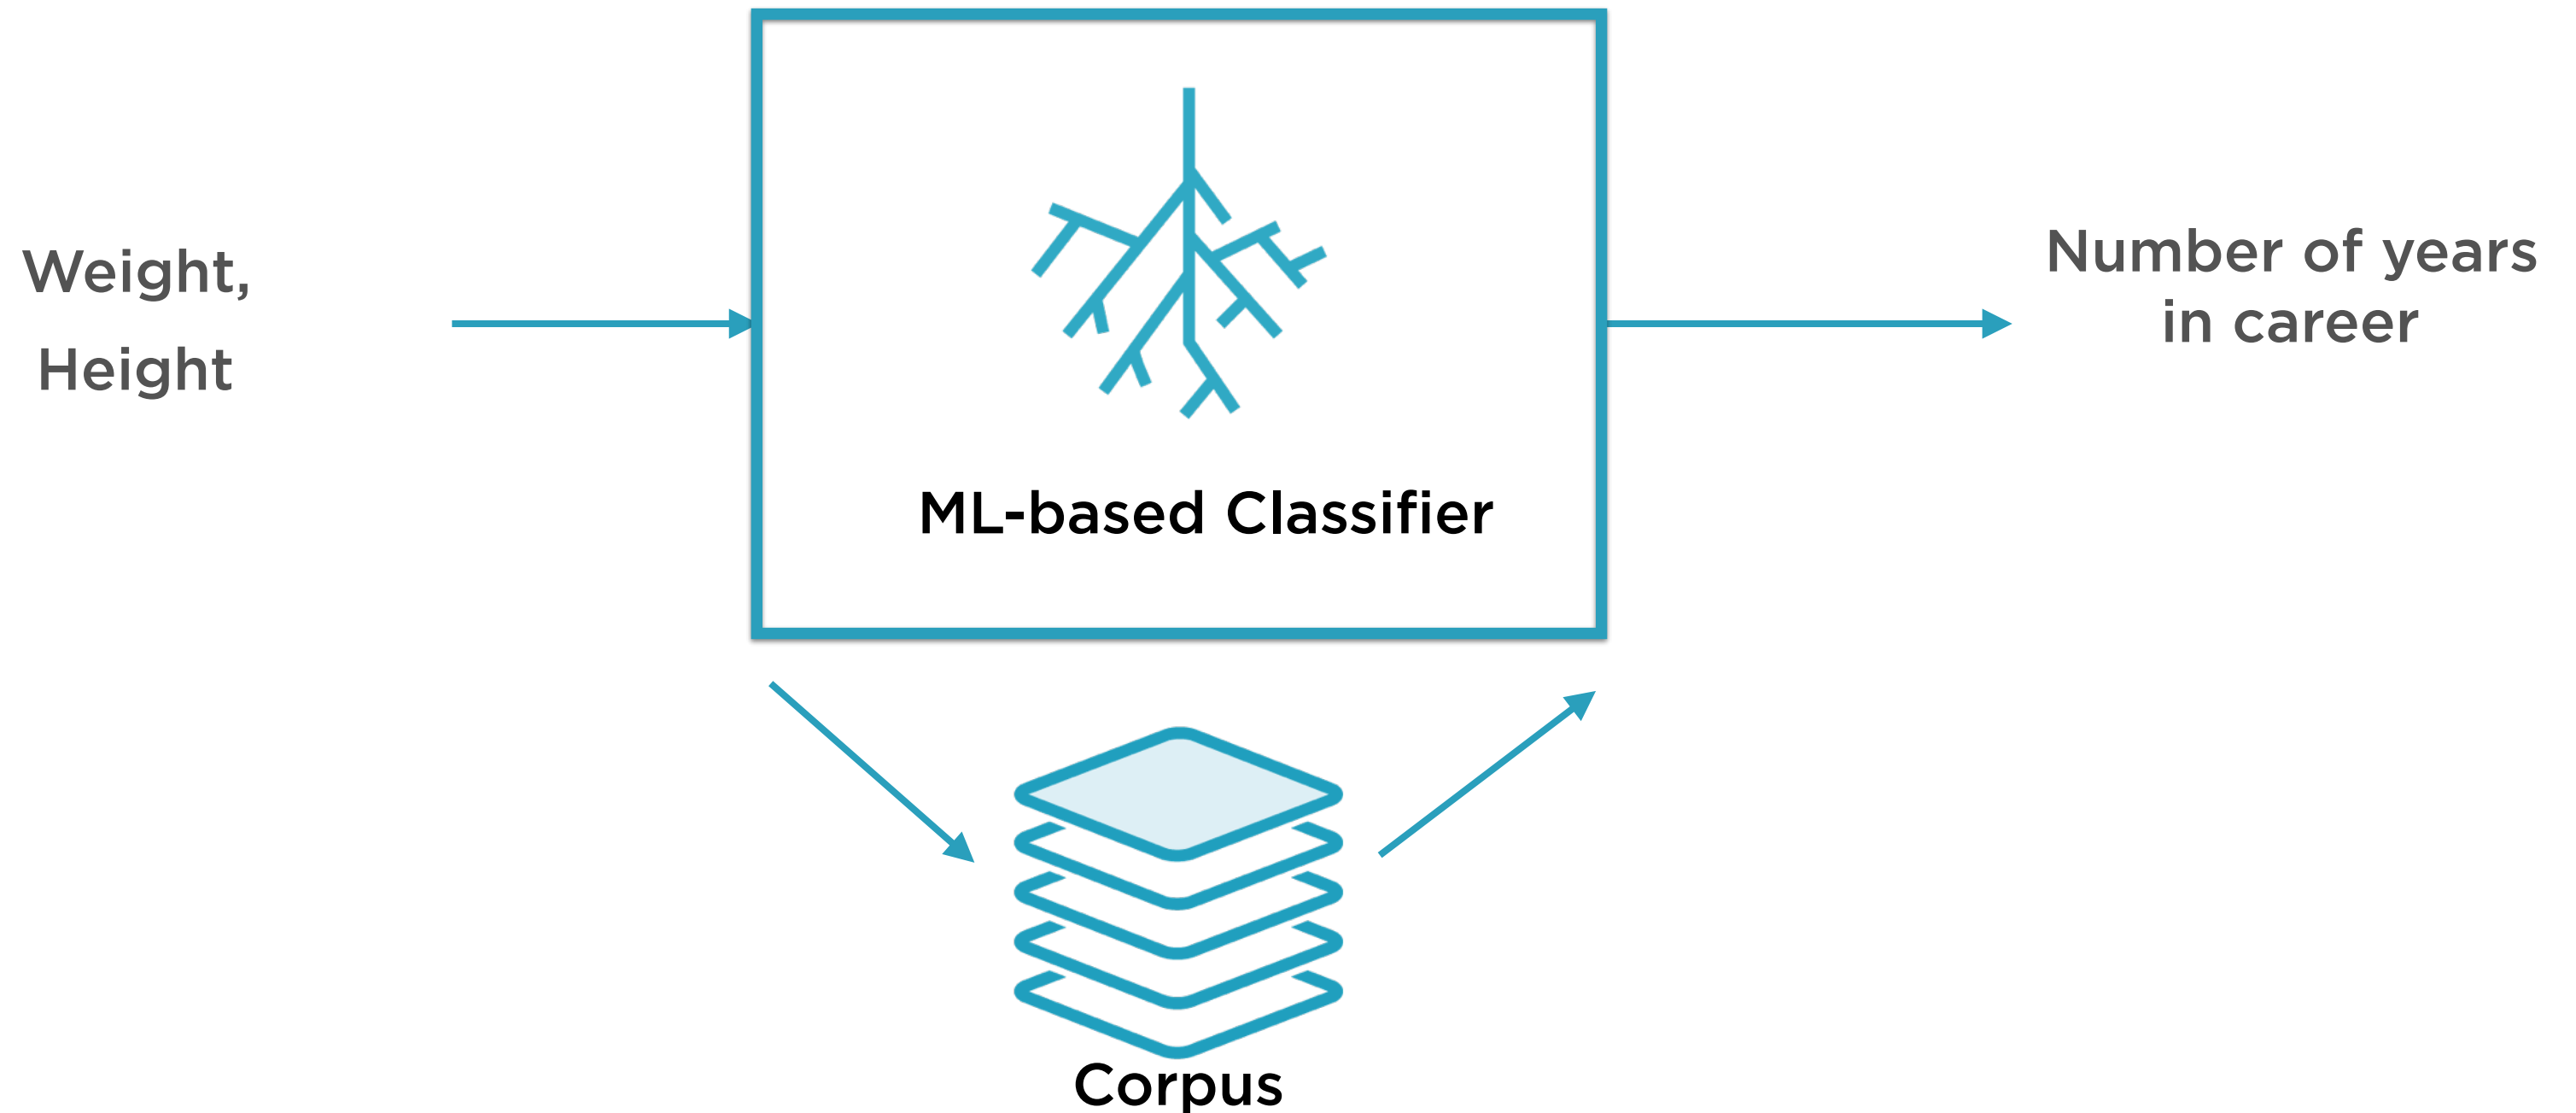
# Decision Trees for Classification




To solve

- Traverse tree to find right node
- Return **most frequent label** of all training data points in that node

# Decision Trees for Regression



# Decision Trees for Regression



To solve

- Traverse tree to find right node
- Return **average number of years** of all training data points in that node

# Muggsy Bogues



**Shortest player ever in the NBA**

**5'3" and 135 lbs**

**Our tree would classify him as Jockey**

**No threshold is perfect!**

# Tree Construction



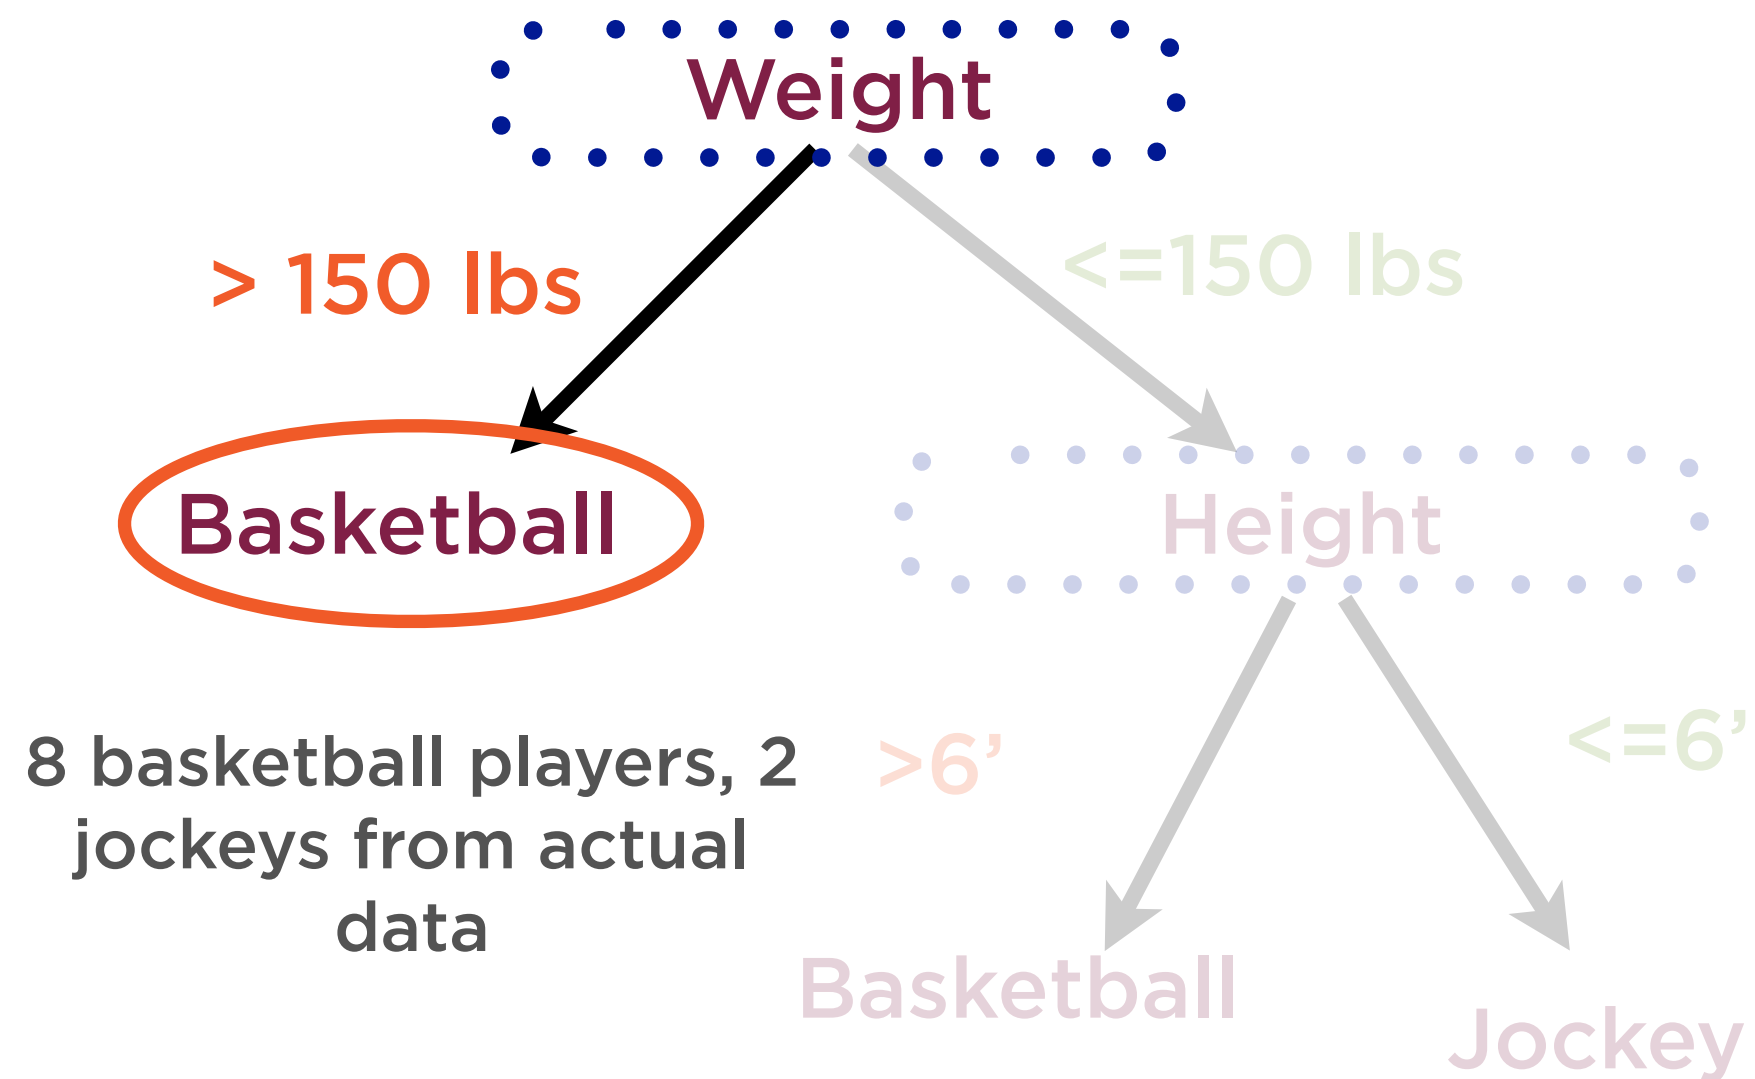
**CART optimizes tree construction**

**Minimizes “impurity” of each node**

**Impurity ~ misclassified data points**



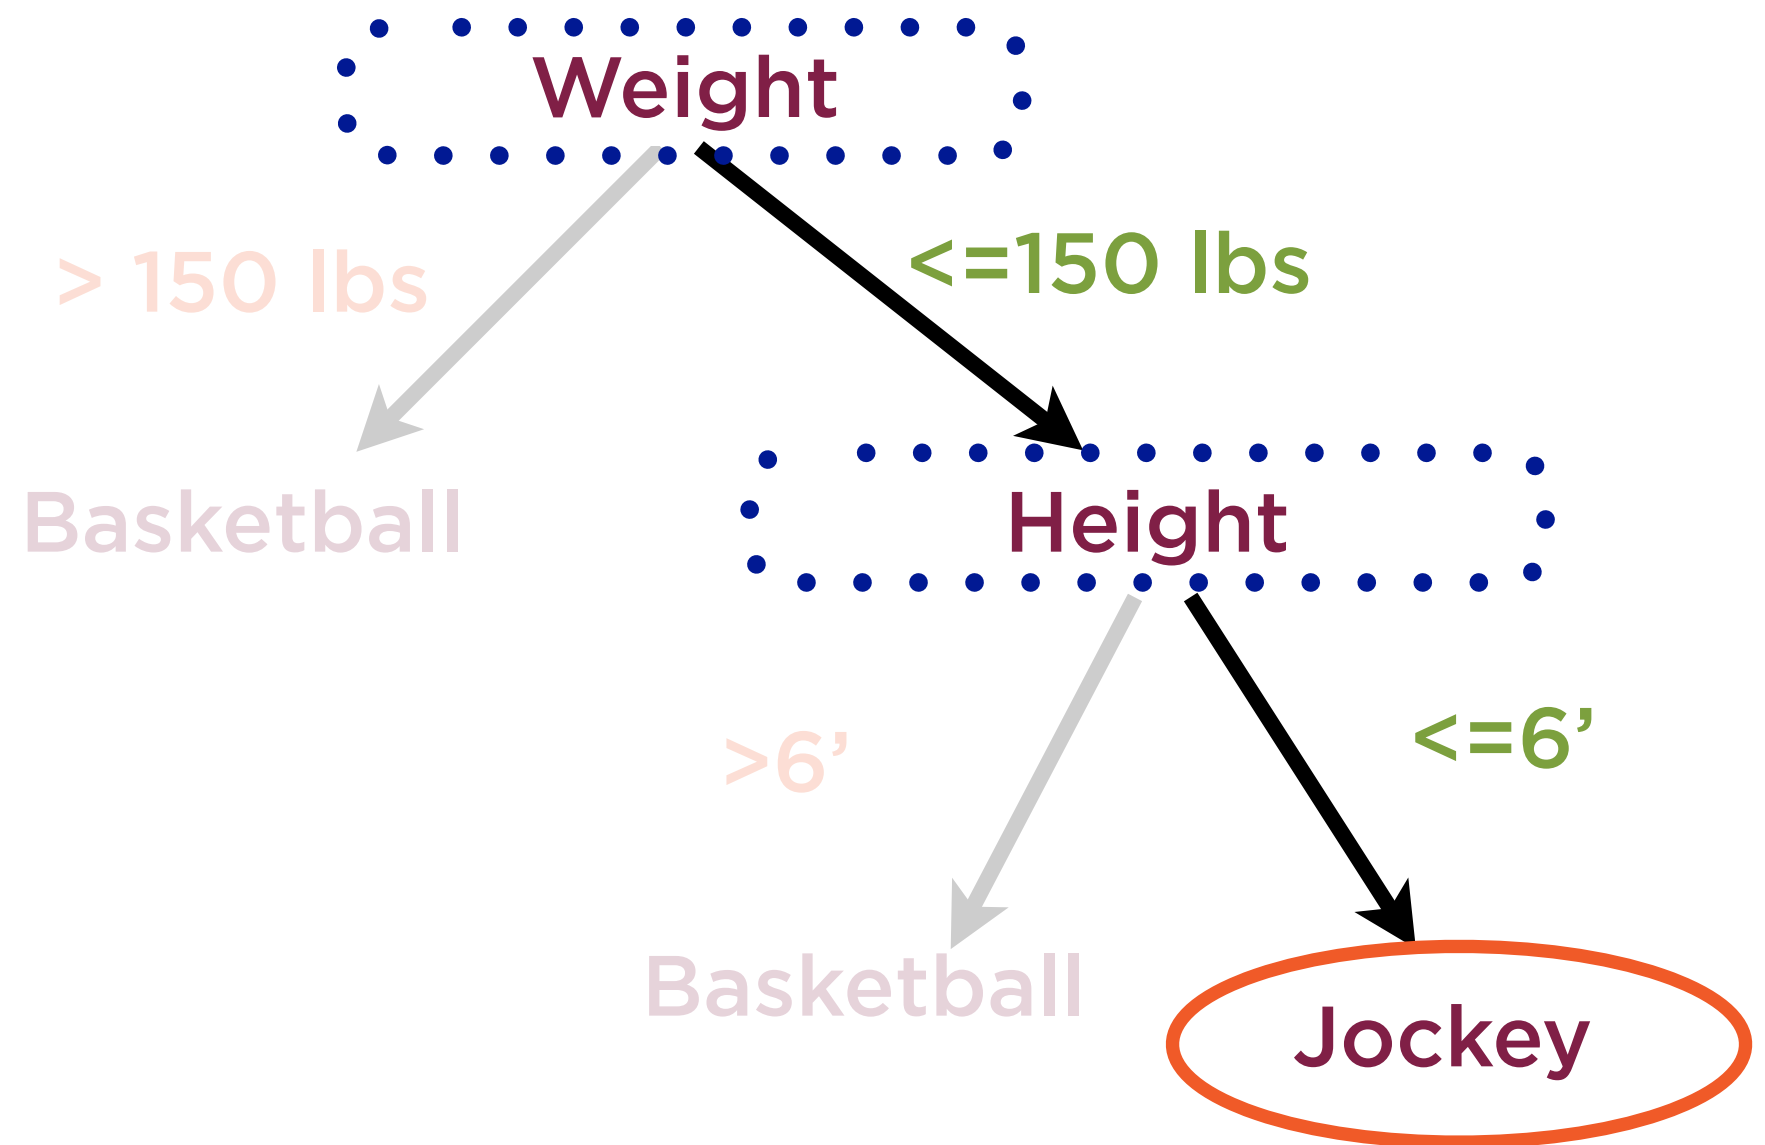
# Impurity







# Impurity




7 jockeys, 3 basketball  
players from actual  
data

Two ways to  
measure impurity


- Gini impurity
- Entropy

Yield similar trees

# Tree Construction




# Gini Impurity



CART seeks to minimize Gini impurity at each node

Gini impurity is found from rule violations in training data

# Gini Impurity




In training data:

**100 samples with height > 6'**

- 95 basketball players
- 5 jockeys

$$G_i = 1 - (95\%)^2 - (5\%)^2 = 0.095$$

# Gini Impurity



In training data:


**100 samples with height  $\leq 6'$**

- 0 basketball players
- 100 jockeys

$$G_i = 1 - (0\%)^2 - (100\%)^2 = 0$$

# Gini Impurity

$$G_i = 0.49875$$




200 samples (sum of the leaf nodes)

- 95 basketball players
- 105 jockeys

$$G_i = 1 - (95/200)^2 - (105/200)^2$$
$$= 0.49875$$

# Advantages of Decision Trees




**“White Box” ML ~ leverage experts**

**Non-parametric**

- **Little hyperparameter tuning**
- **Little data prep**

# Drawbacks of Decision Trees



## Prone to overfitting

- Common risk with non-parametric

## Unstable

- Small changes in data cause big changes in model

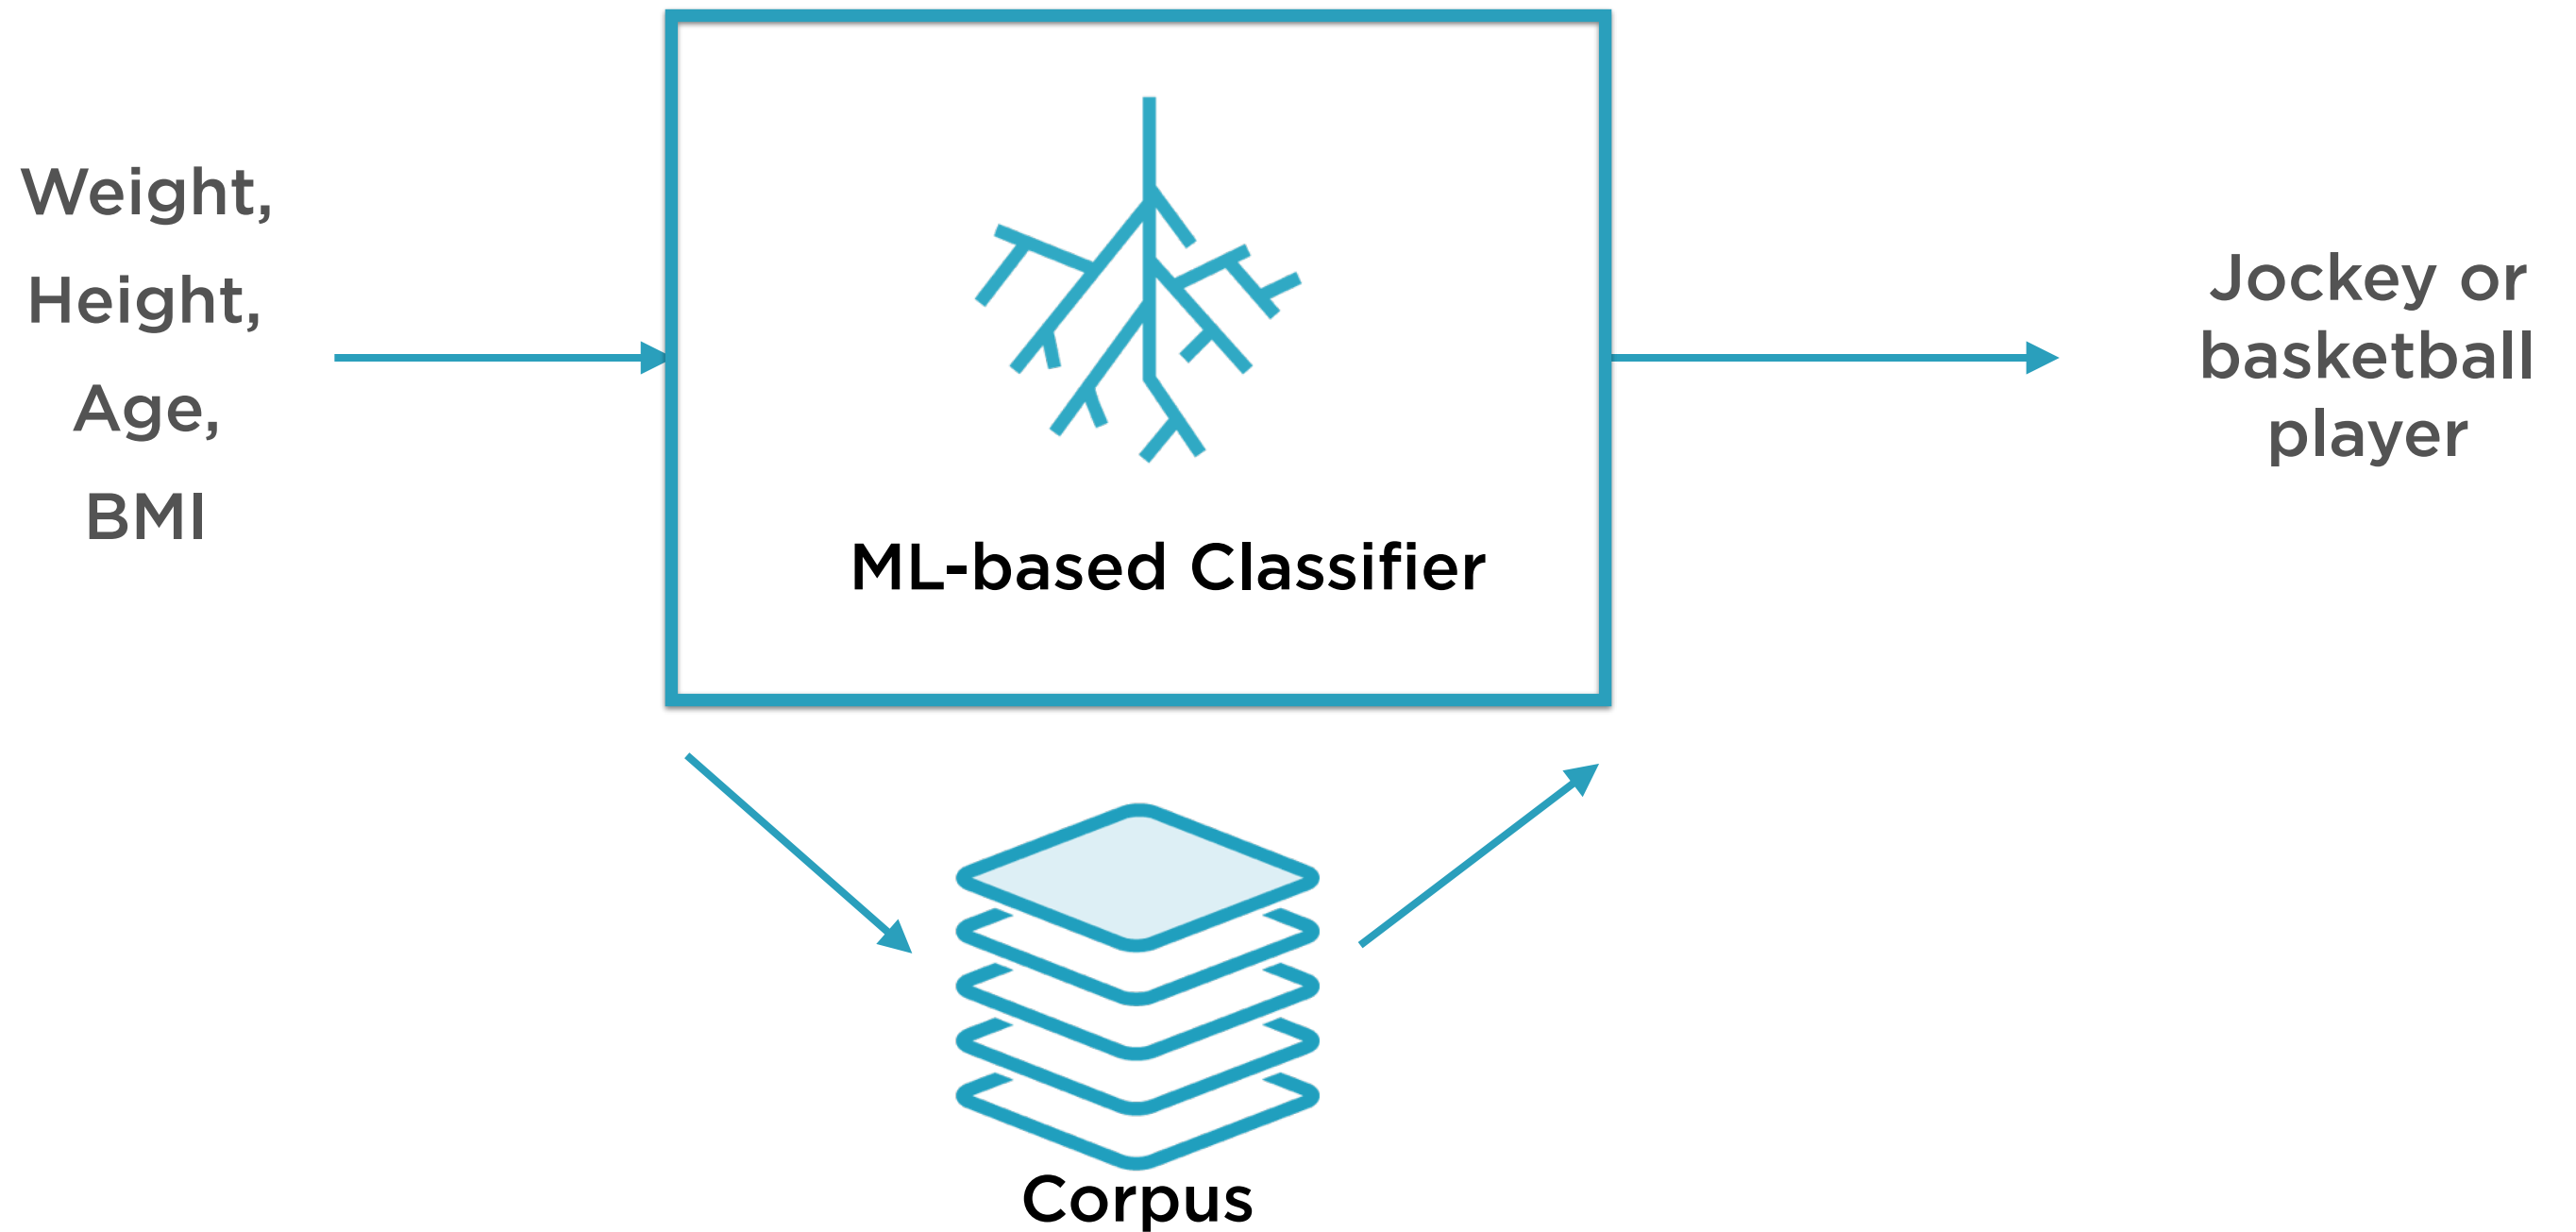


# Demo

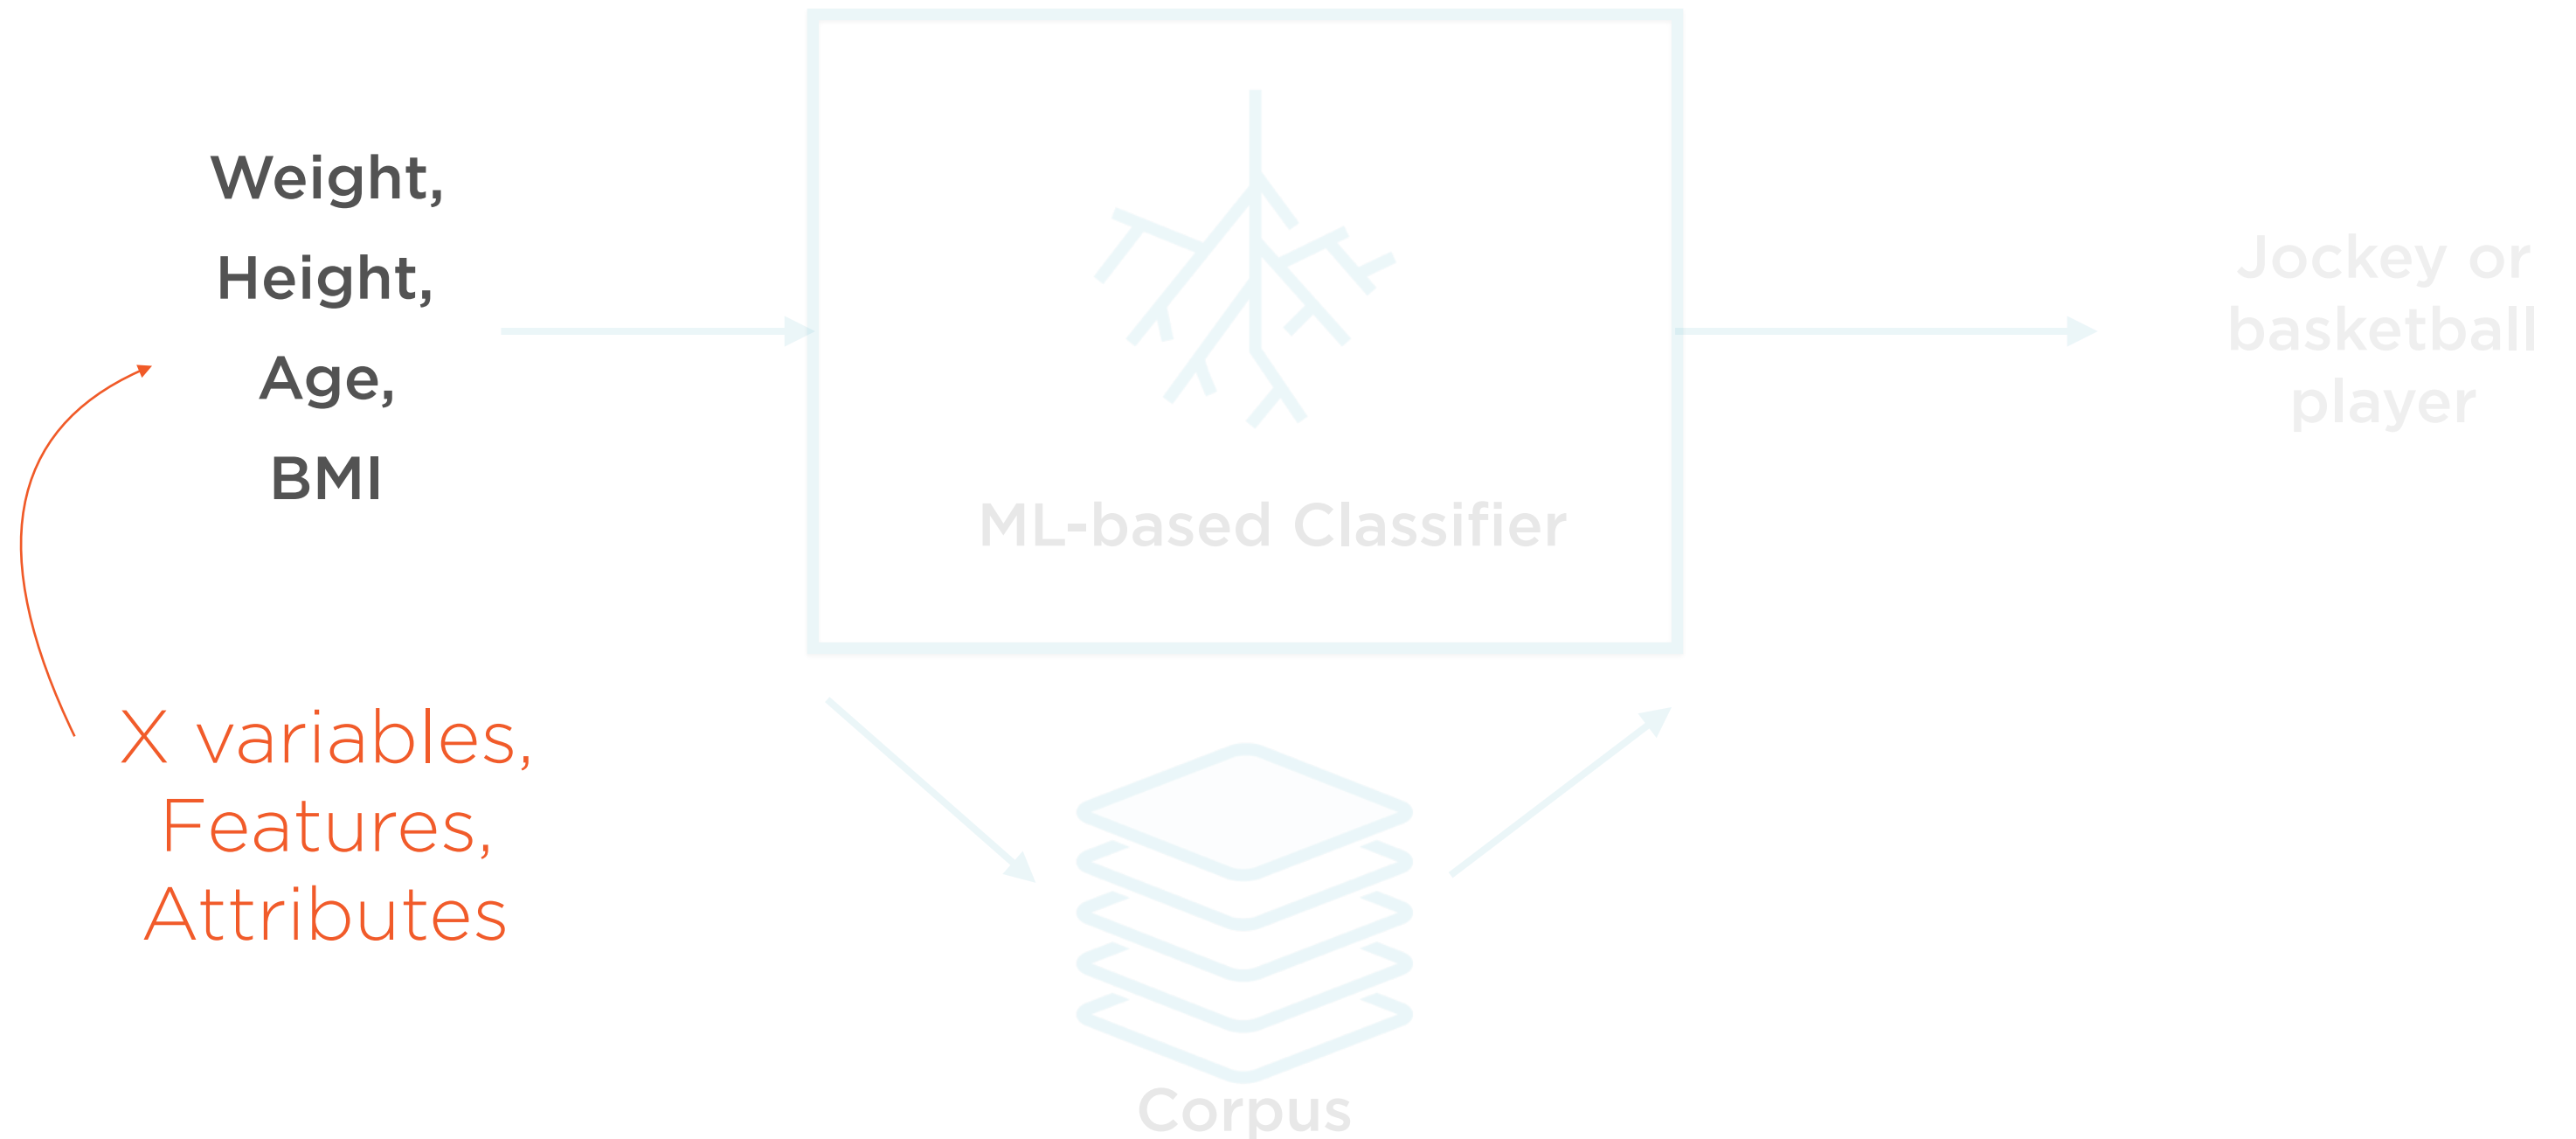
**Implement classification using decision trees in spark.mllib**

**Data in the CSV as well as the LIBSVM format**


# LIBSVM Data Format




# LIBSVM Data Format




# LIBSVM Data Format




# LIBSVM Data Format




# LIBSVM Data Format



# LIBSVM Data Format

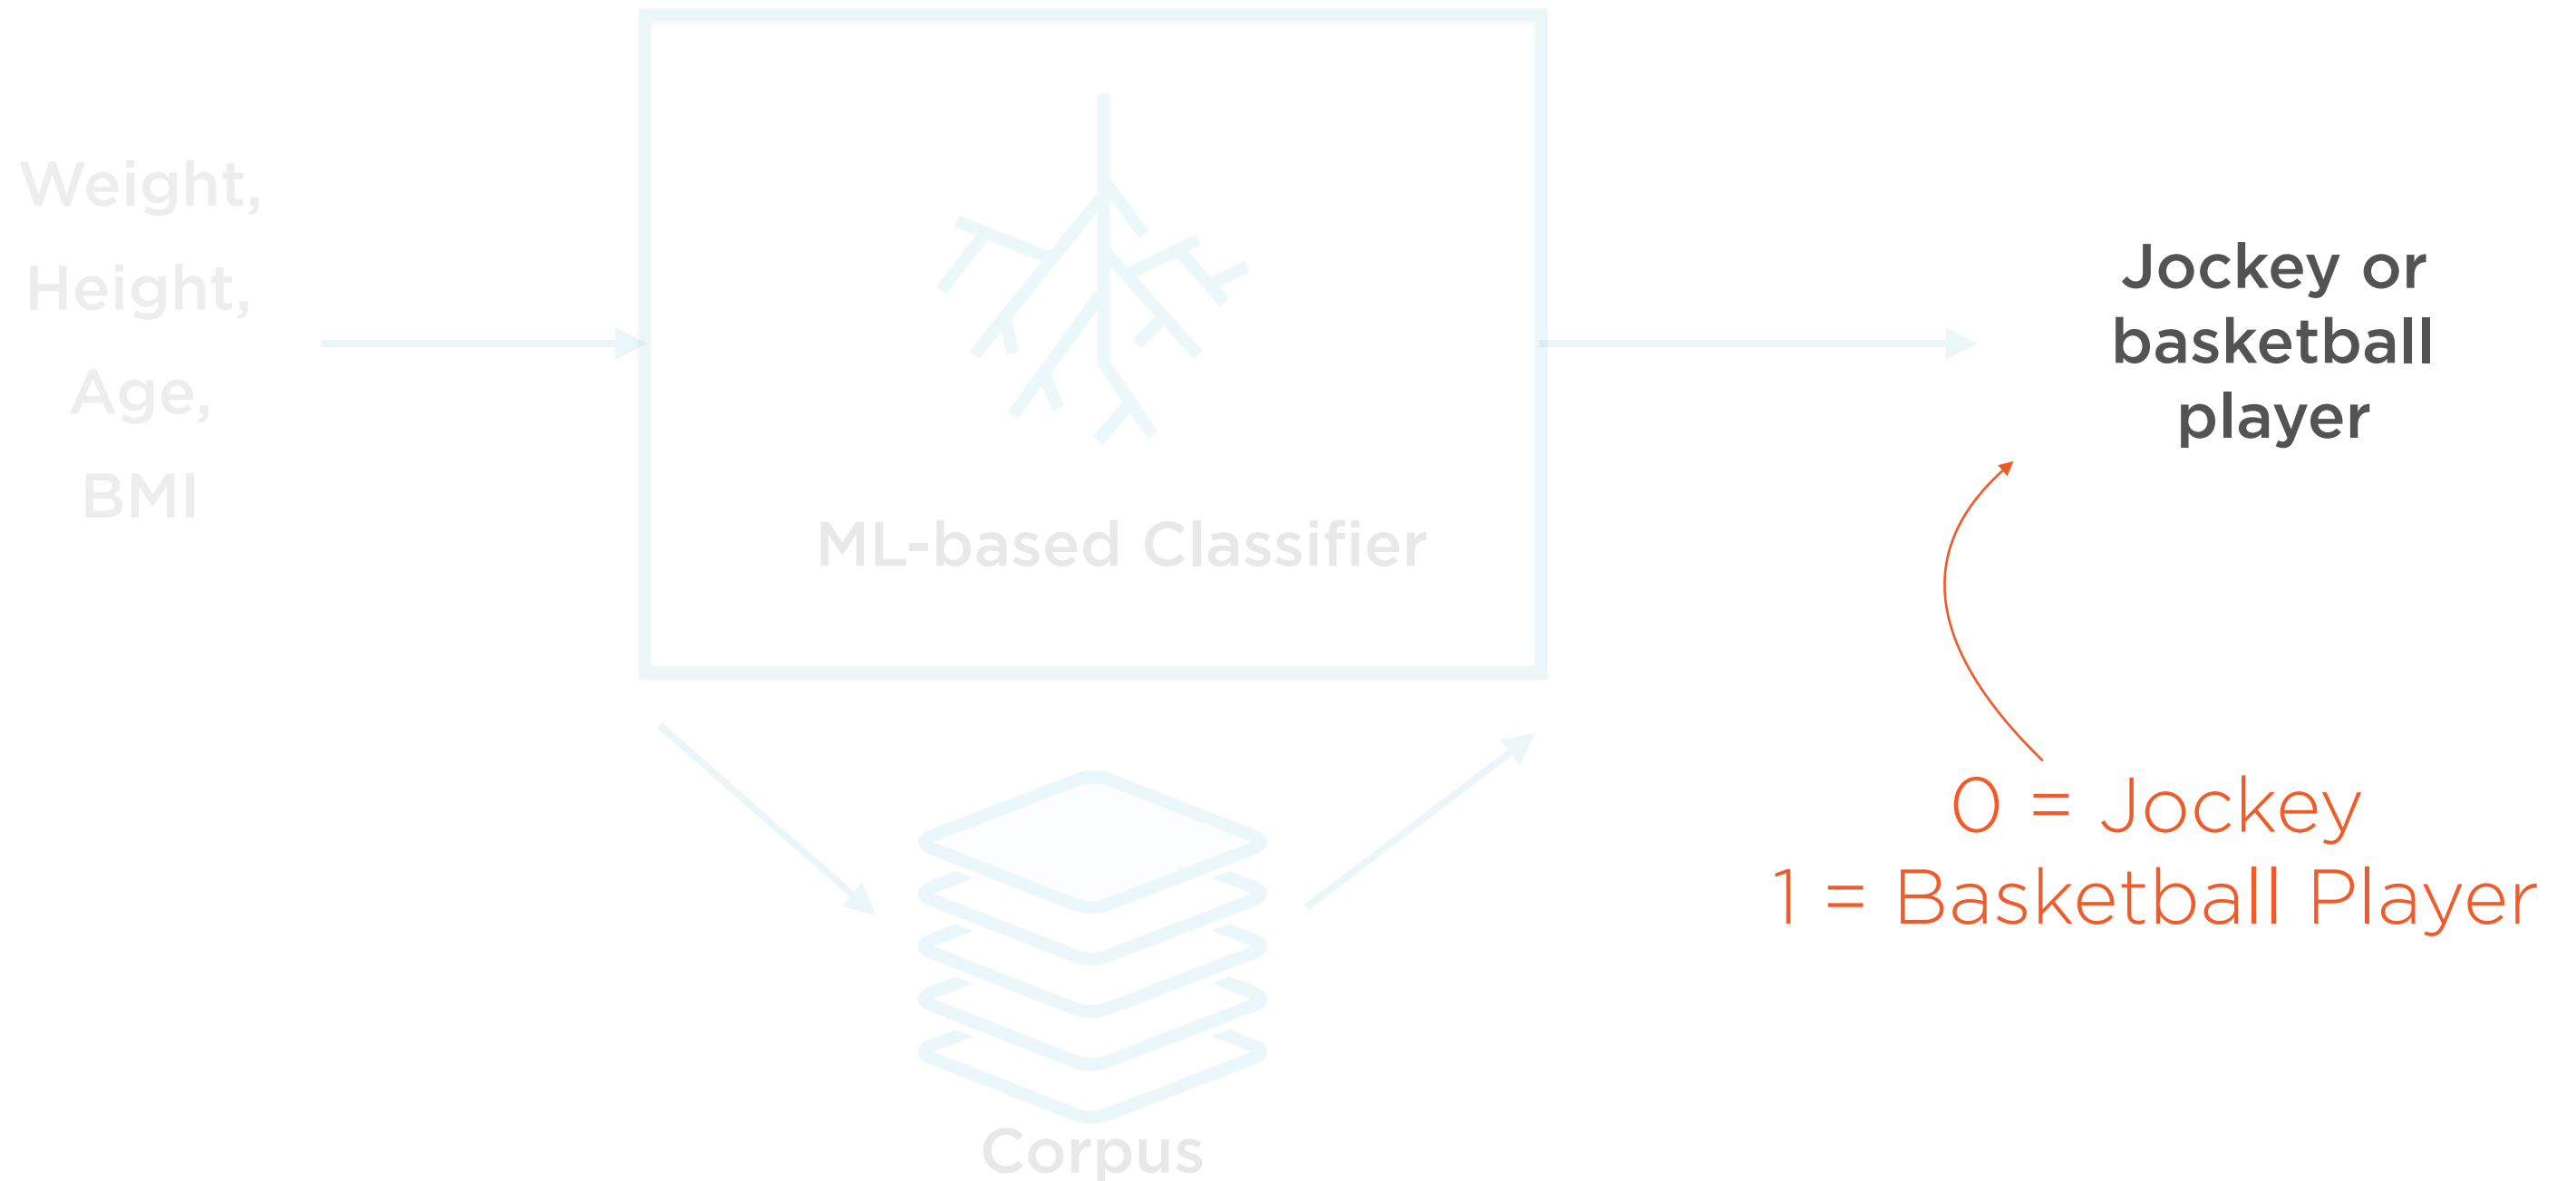


# LIBSVM Data Format





# LIBSVM Data Format



# LIBSVM Data Format

**<label> <index1>:<value1> <index2>:<value2>...**

**0 1:230 2:188 3:32 4:29.4**

1 line per instance

Each line ends with '\n'

Sparse - missing attributes can be omitted

# LIBSVM Data Format

**<label>** <index1>:<value1> <index2>:<value2>...

0 1:230 2:188 3:32 4:29.4

0 = Jockey

# LIBSVM Data Format

<label> **<index1>**:<value1> <index2>:<value2>...

0 **1:230** 2:188 3:32 4:29.4

Index = 1 for attribute **Weight**

# LIBSVM Data Format

<label> <index1>:<value1> <index2>:<value2>...

0 1:230 2:188 3:32 4:29.4

Weight in lbs = 230

# LIBSVM Data Format

<label> <index1>:<value1> **<index2>:<value2>**...

0 1:230 **2:188** 3:32 4:29.4

Index = 2 for attribute **Height**

# LIBSVM Data Format

<label> <index1>:<value1> **<index2>:<value2>...**

0 1:230 **2:188** 3:32 4:29.4

Height in cm = 188

# LIBSVM Data Format

<label> <index1>:<value1> <index2>:<value2>...

0 1:230 2:188 **3:32** 4:29.4

Index = 3 for attribute **Age**  
Age in years = 32



# LIBSVM Data Format

<label> <index1>:<value1> <index2>:<value2>...

0 1:230 2:188 3:32 **4:29.4**

Index = 4 for attribute **BMI**  
BMI ratio = 29.4

# LIBSVM Data Format

**<label> <index1>:<value1> <index2>:<value2>...**

**0 1:230 2:188 3:32 4:29.4**

1 line per instance

Each line ends with '\n'

Sparse - missing attributes take value 0

# Missing Attributes

**<label> <index1>:<value1> <index2>:<value2>...**

**0 1:145 2:158 3:39**

Value for index 4 is missing

No worries - can calculate from height and weight

Sparse - missing attributes can be omitted

# Summary

**Spark 1.x provided powerful support for ML in spark.mllib**

**Spark 2.0 goes further with spark.ml**

**Faster execution**

**Ease of hyperparameter tuning**

**ETL support with ML pipelines**

**spark.mllib currently has more features but will be deprecated in the future**